

Autoregressive Models

MUSTAFA HAJIJ

Introduction

Autoregressive model over x factor the joint distribution as the following product of conditionals:

$$p(\mathbf{x}) = p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

The problem becomes modeling the conditional probability distribution : $P(x_i | x_1, x_2, \dots, x_{i-1})$

Introduction

Autoregressive model over x factor the joint distribution as the following product of conditionals:

$$p(\mathbf{x}) = p(x_1, \dots, x_n) = \prod_{i=1}^n p(x_i | x_1, \dots, x_{i-1})$$

The problem becomes modeling the conditional probability distribution : $P(x_i | x_1, x_2, \dots, x_{i-1})$

Interpretation : Given previous values x_1, x_2, \dots, x_{i-1} , these models do not output a value for x_i , they output the predictive probability distribution $P(x_i | x_1, x_2, \dots, x_{i-1})$ for x_i .

Introduction

Any autoregressive model can be run sequentially to generate a new sequence : start with your seed x_1, x_2, \dots, x_k and predict x_{k+1} . Then use x_2, x_3, \dots, x_{k+1} to predict x_{k+2} , and so on.

PixelCNN

- **PixelCNN** is a generative model that predicts the next pixel in an image given the previous pixels.
- It leverages the idea of autoregressive modeling, where the image is generated pixel by pixel in a specific order.

PixelCNN

- **Architecture Overview:**

- Consists of multiple convolutional layers.
- Each layer is designed to ensure that the generation of a pixel depends only on previously generated pixels.

PixelCNN

- **Masking Concept:** Masks are used in the convolutional layers to prevent information from flowing from future pixels to the current pixel.
- **Types of Masks:**
 - **A-type Mask:** Used in the first layer to ensure that a pixel does not depend on itself.
 - **B-type Mask:** Used in subsequent layers, allowing the pixel to depend on itself and previous pixels.





Figure 5-13. Left - a convolutional filter mask. Type A masks the central pixel and Type B does not mask the central pixel. Right - a mask applied to a set of pixels to predict the distribution of the central pixel value (source: Conditional Image Generation with PixelCNN Decoders, van den Oord et al. <https://arxiv.org/pdf/1606.05328>).

PixelCNN

- **Problems: very slow to sample, for an image of 64 by 64, we need to sample 64 by 64!**

One solution is to replace the head of softmax over 256 by a mixture distribution.
For instance in the following example, we have a mixture with three kernels.



PixelCNN

occluded


completions

original



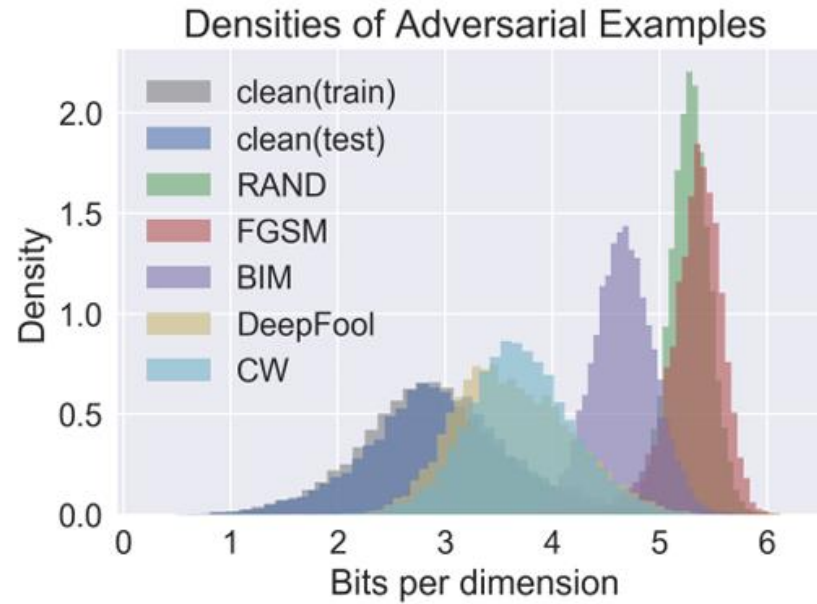
Applications : Adversarial Attacks and Anomaly detection

Machine learning models are vulnerable to adversarial examples



Can we detect such examples?

Applications : Adversarial Attacks and Anomaly detection



Train a generative model $p(x)$ on clean inputs (PixelCNN)

Given a new input x' , evaluate $p(x')$

Adversarial examples are significantly less likely under $p(x)$

Summary

- Easy to sample from
 - ① Sample $\bar{x}_0 \sim p(x_0)$
 - ② Sample $\bar{x}_1 \sim p(x_1 \mid x_0 = \bar{x}_0)$
 - ③ ...
- Easy to compute probability $p(x = \bar{x})$
 - ① Compute $p(x_0 = \bar{x}_0)$
 - ② Compute $p(x_1 = \bar{x}_1 \mid x_0 = \bar{x}_0)$
 - ③ Multiply together (sum their logarithms)
 - ④ ...
 - ⑤ Ideally, can compute all these terms in parallel for fast training