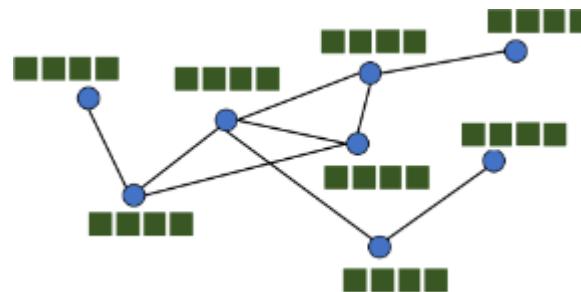


# Graph Generative Models

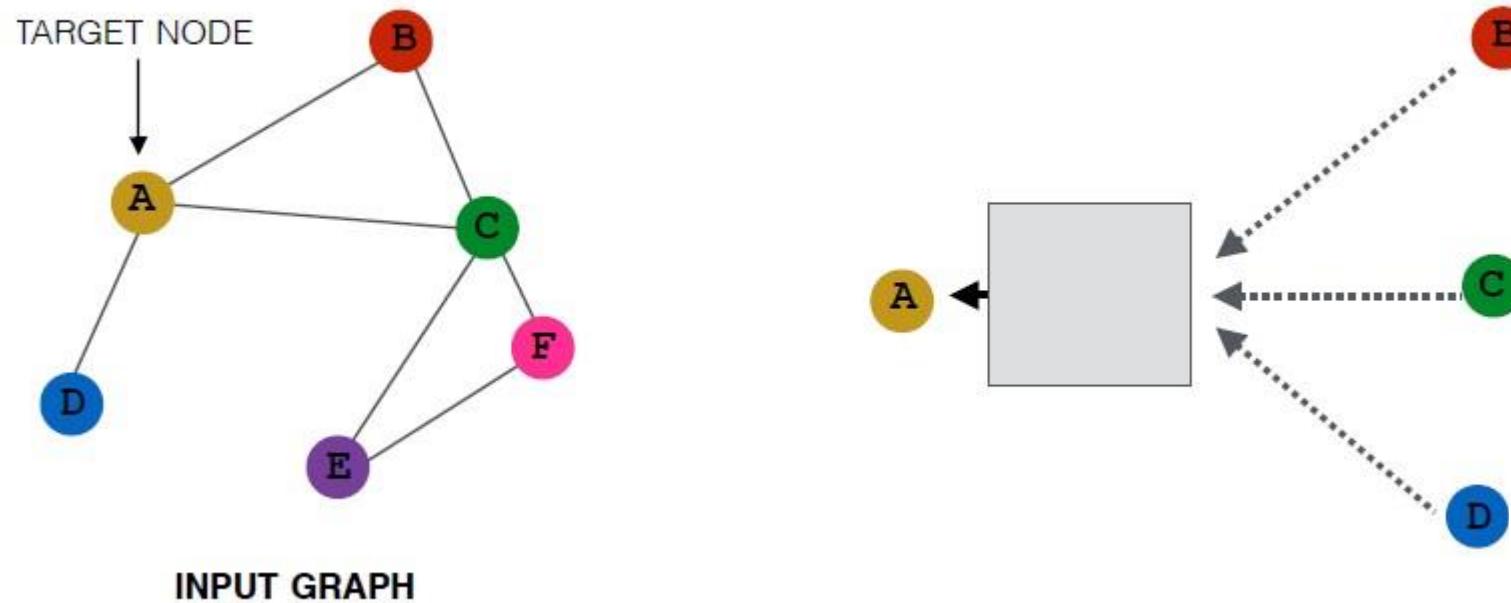
## Problem setup

Assume we have a graph  $G$ :

- $V$  is the vertex set
- $A$  is the adjacency matrix (assume binary)
- $X \in \mathbb{R}^{m \times |V|}$  is a matrix of node features
- $v$ : a node in  $V$ ;  $N(v)$  : the set of neighbors of  $v$ .



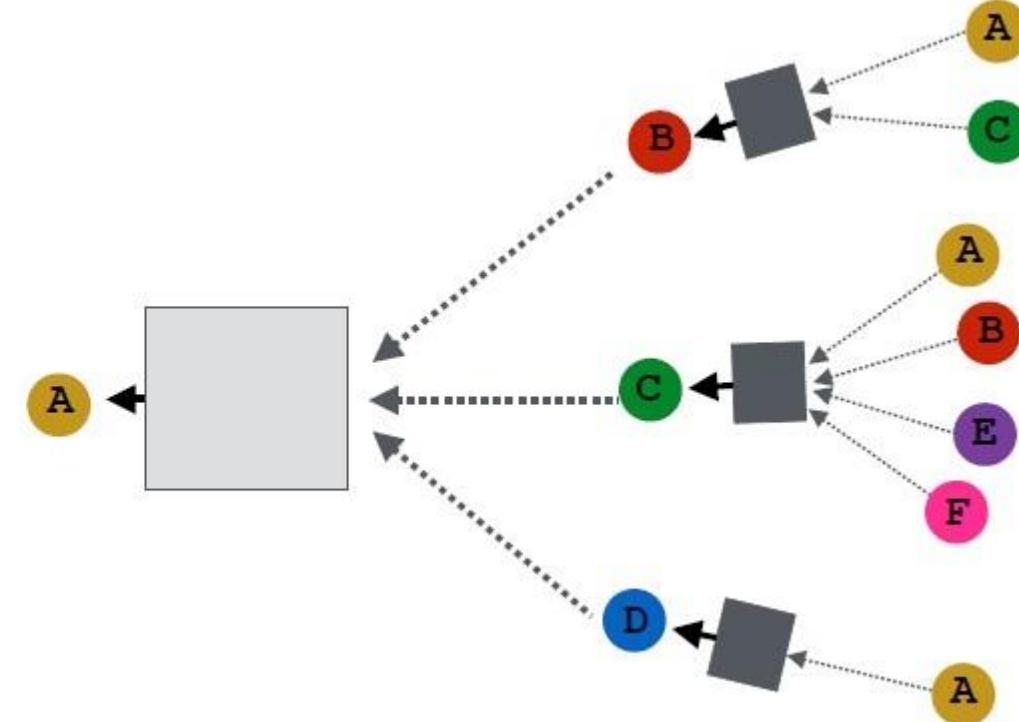
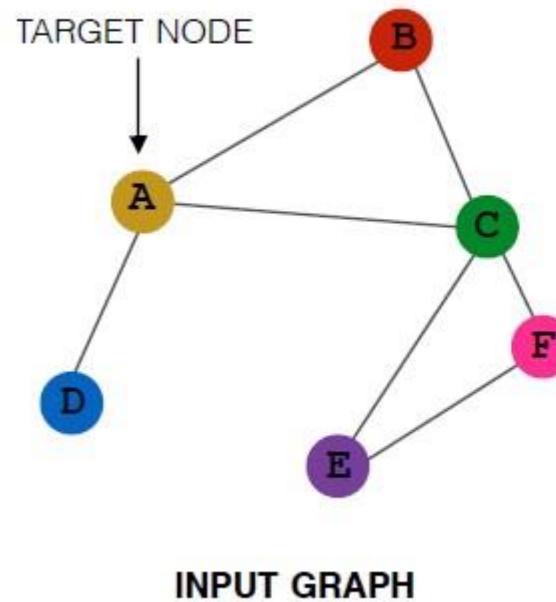
## Graph Neural Networks



**Key idea:** Each node defines a computation graph

- Each edge in this graph is a transformation/aggregation function

## Graph Neural Networks



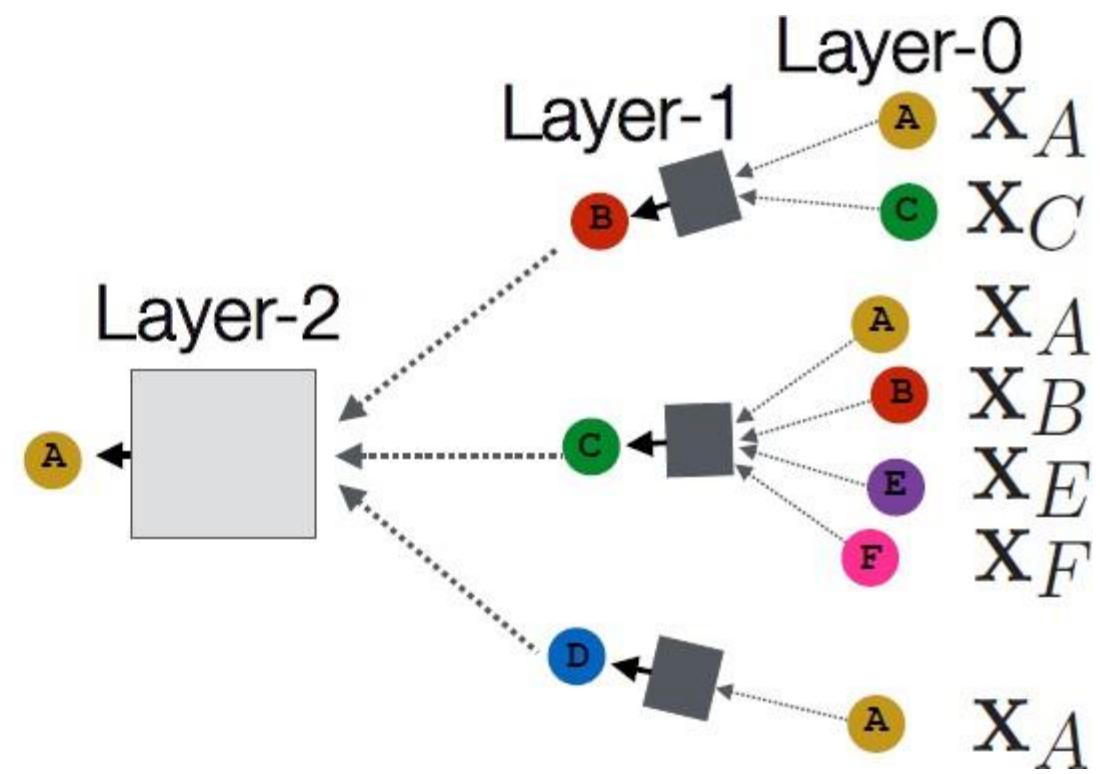
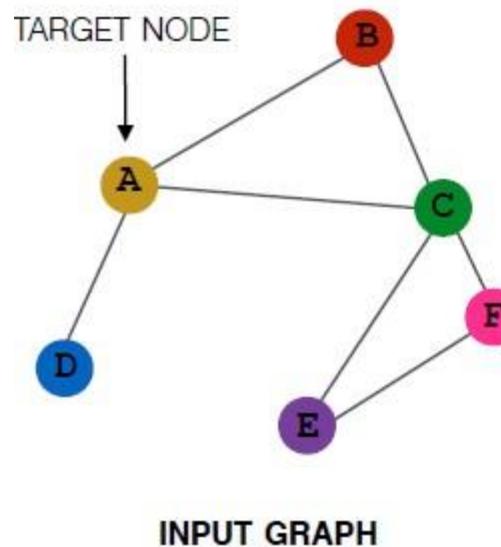
**Key idea:** Each node defines a computation graph

- Each edge in this graph is a transformation/aggregation function

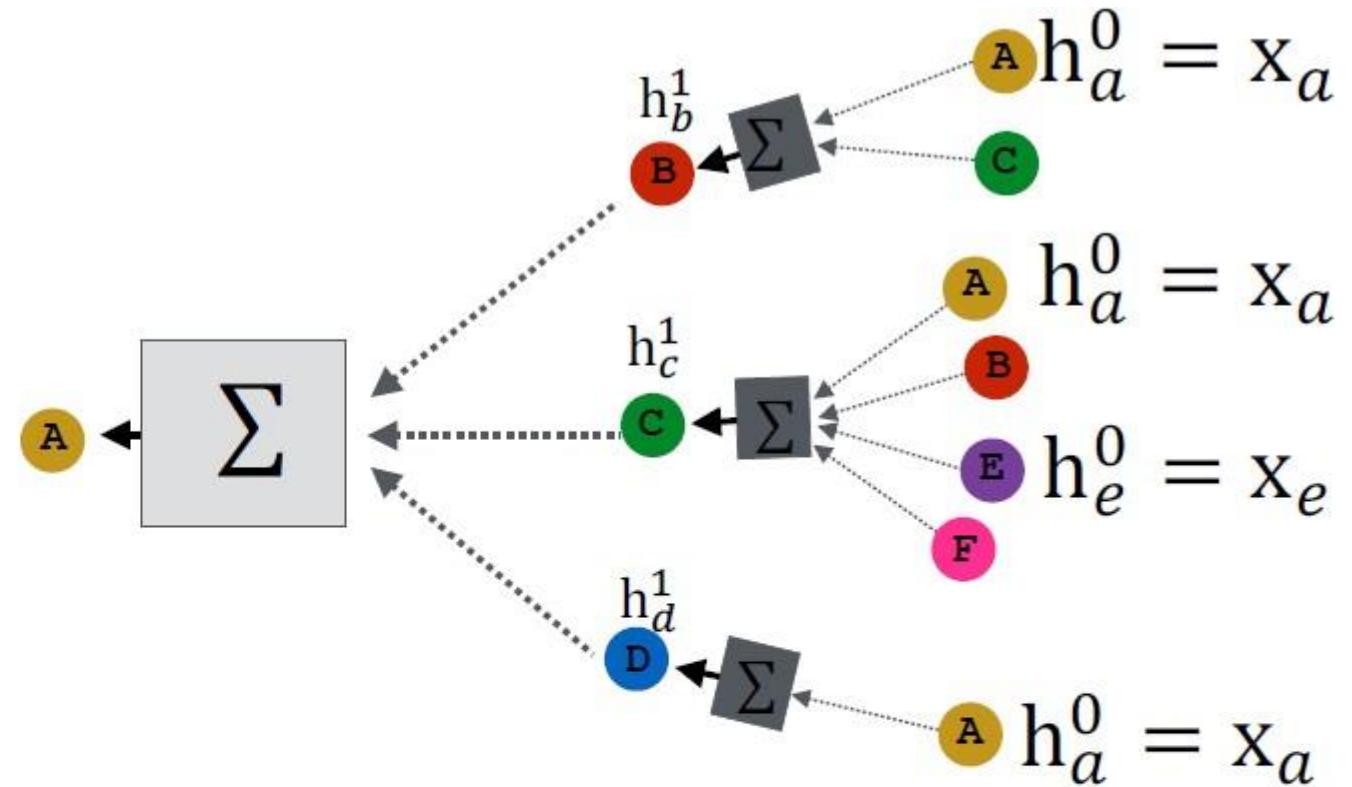
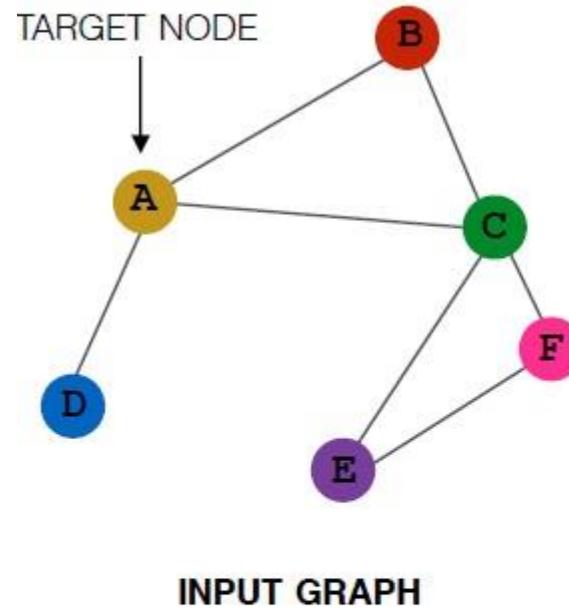
## Graph Neural Networks

Nodes have embeddings at each layer

- Layer-0 embedding of node  $u$  is its input feature,  $x_u$
- Layer- $k$  embedding gets information from nodes that are  $K$  hops away

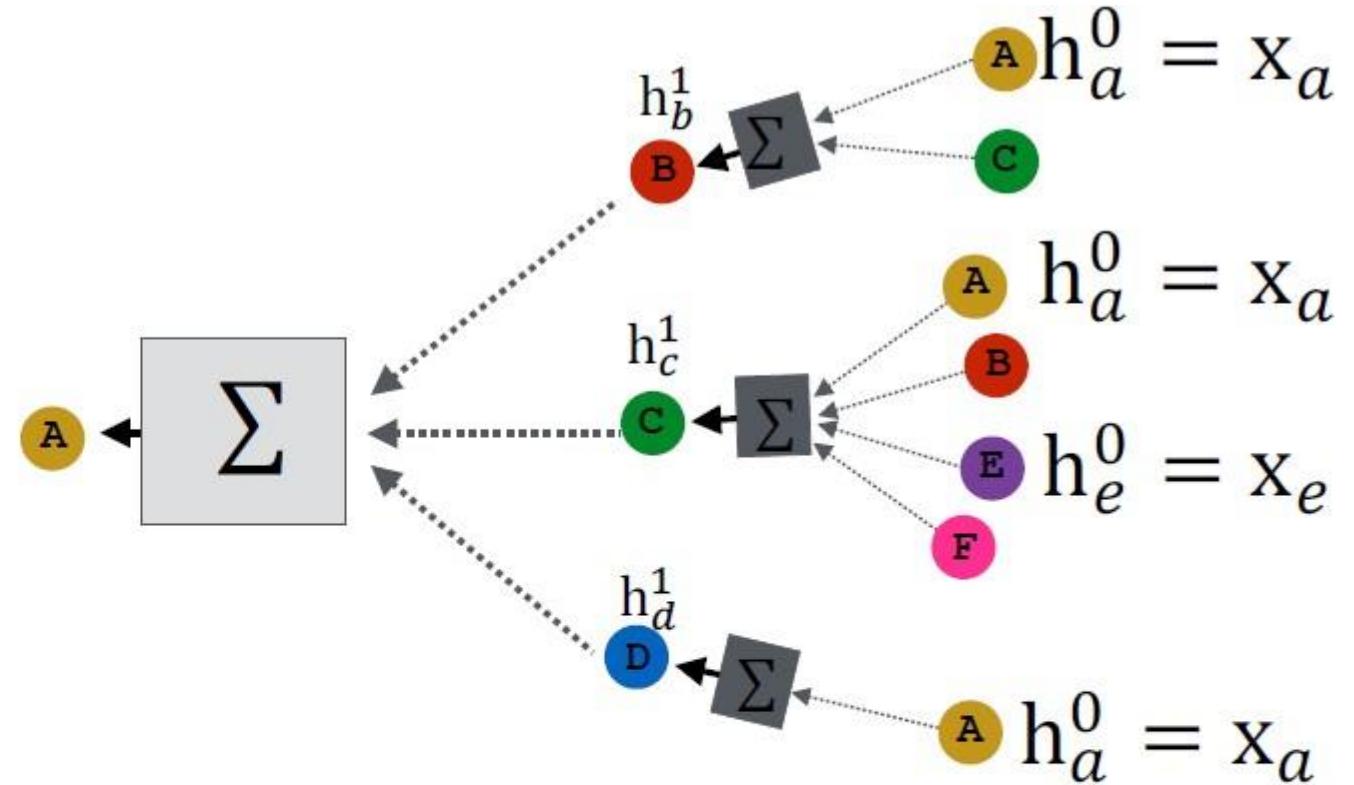
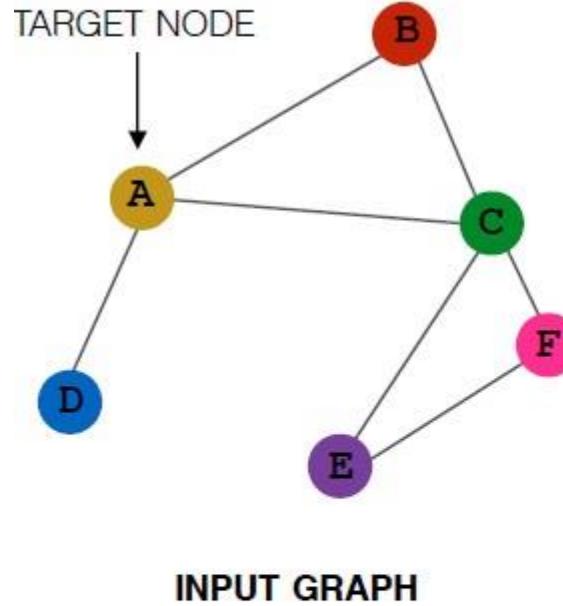


## Graph Neural Networks



$$h_v^{(l+1)} = \sigma(W_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)})$$

## Graph Neural Networks



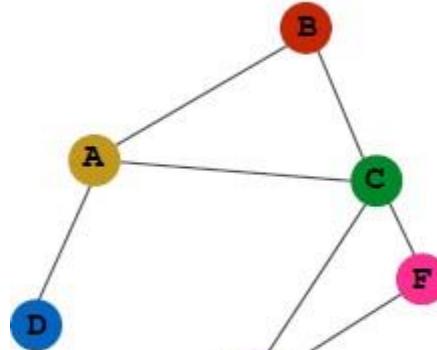
$$h_v^{(l+1)} = \sigma(W_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)})$$

Learnable parameters

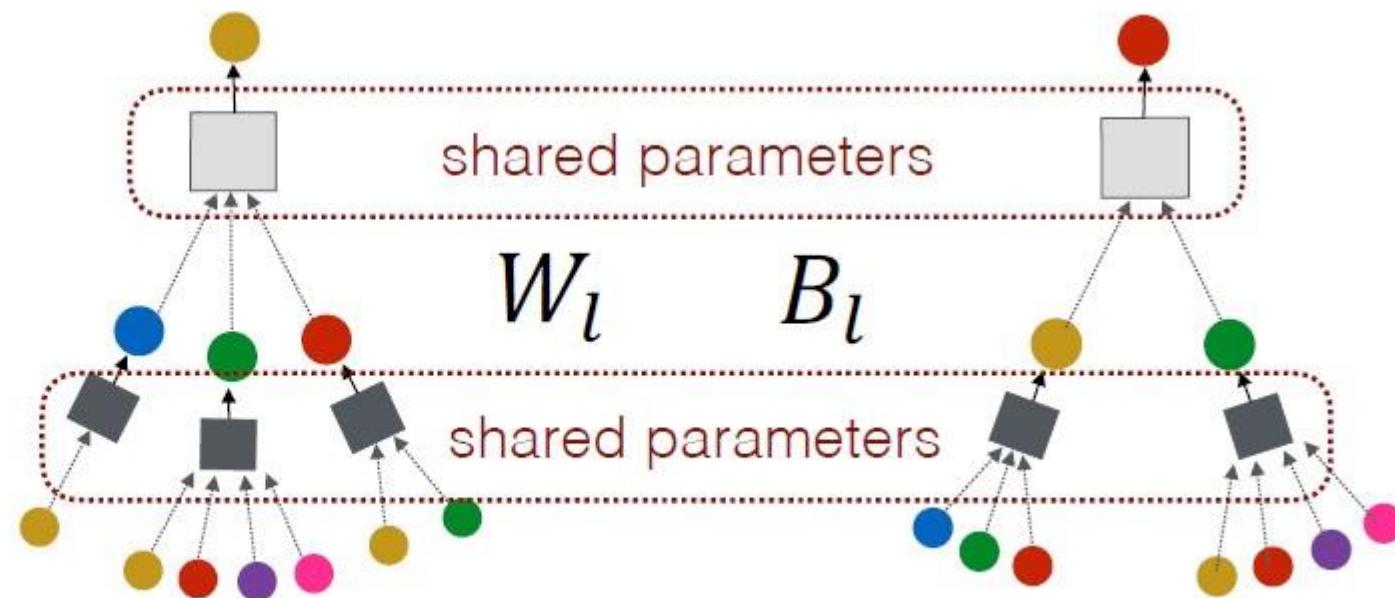
[Image source](#)

## Graph Neural Networks

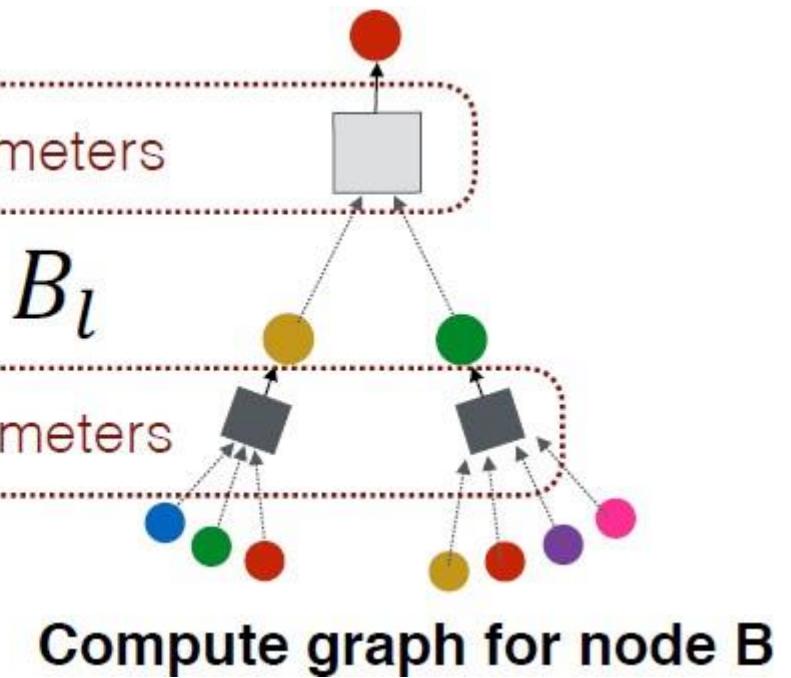
The *same* aggregation parameters are shared for all nodes



INPUT GRAPH

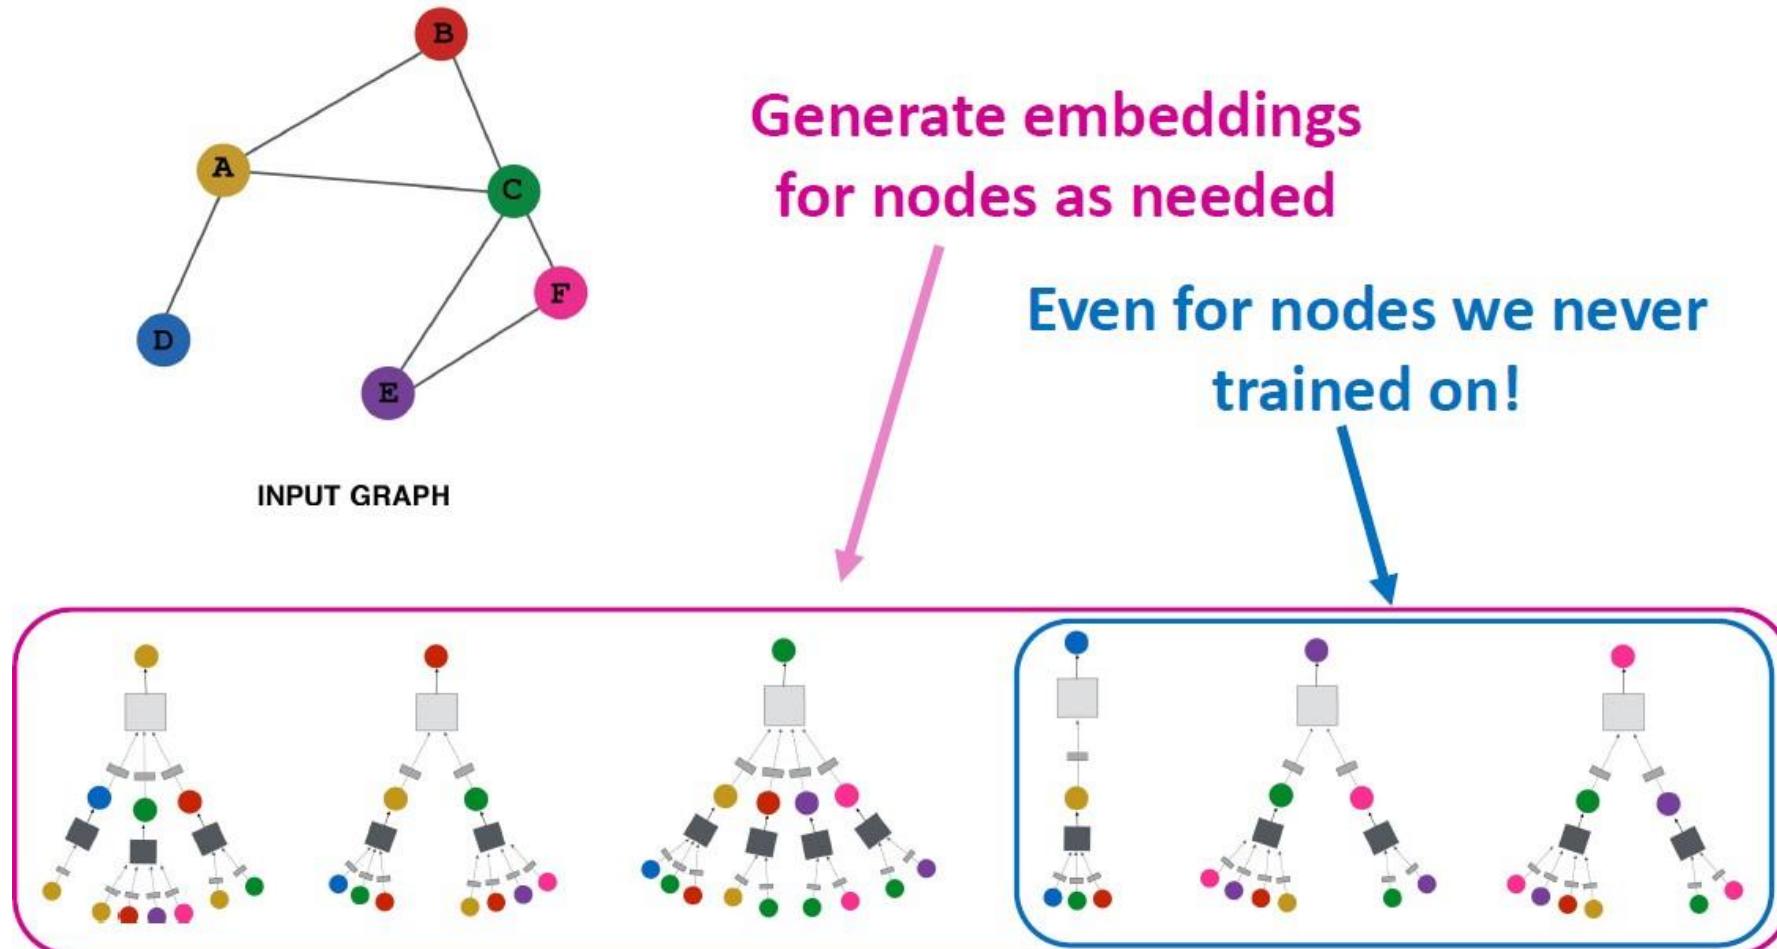


Compute graph for node A



Compute graph for node B

## Graph Neural Networks

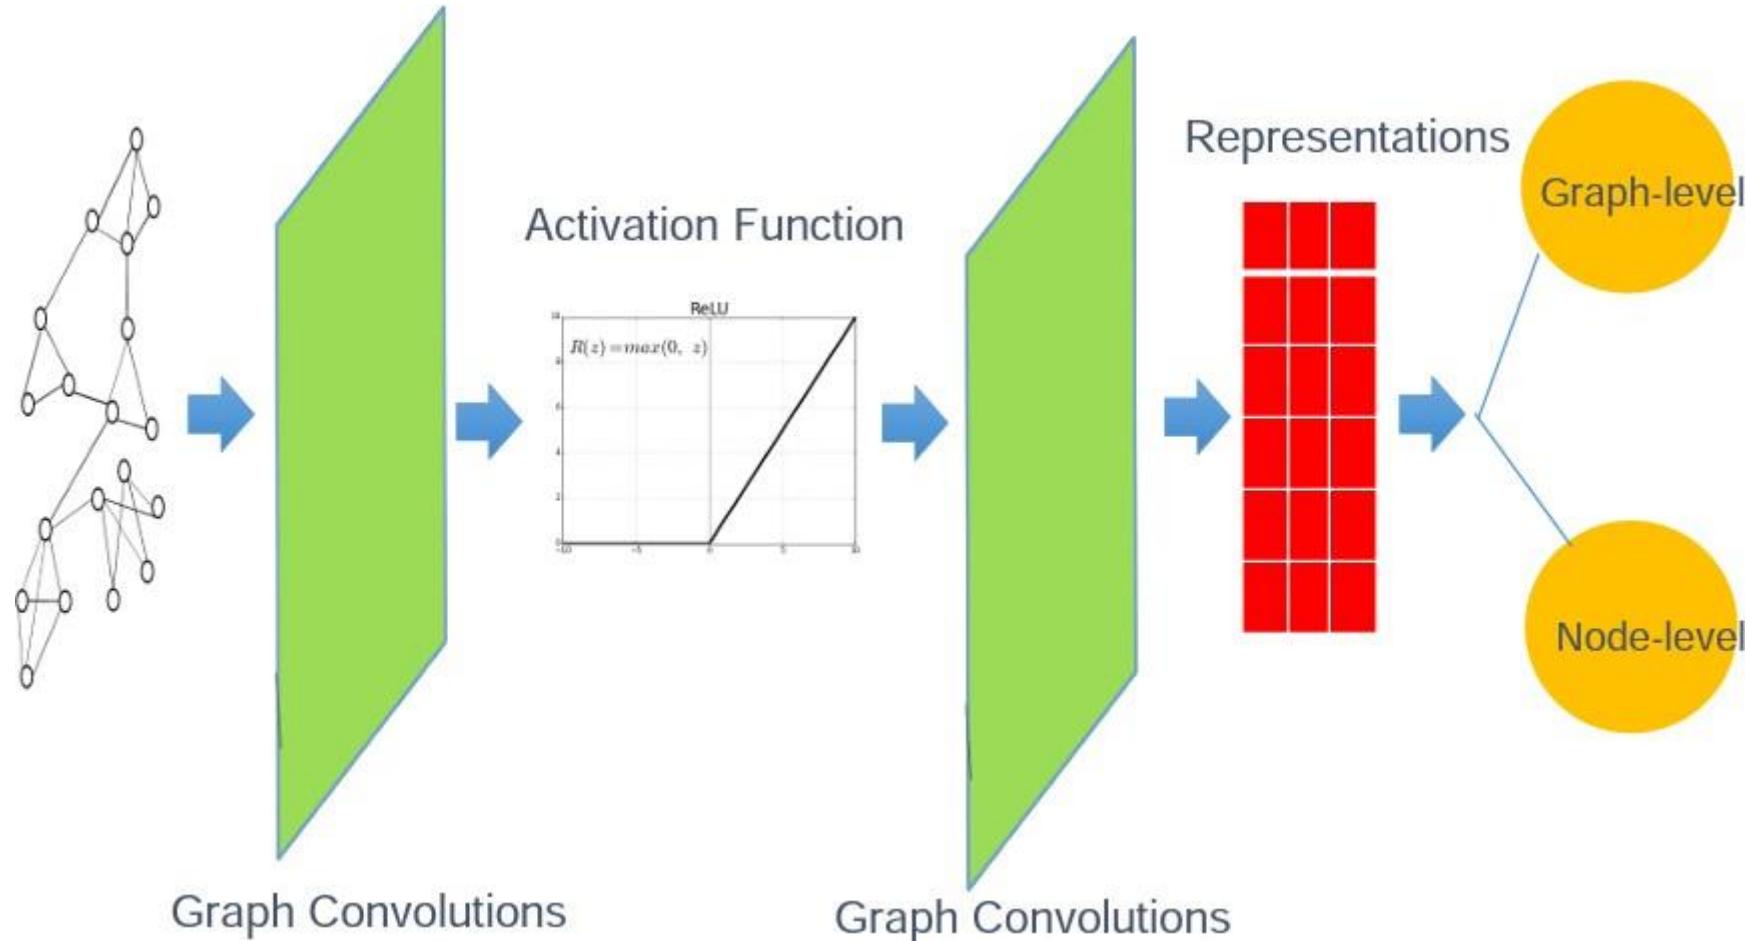


## Graph Convolutional Network

$$h_v^{k+1} = \sigma(W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|} + B_k h_v^k)$$

**Key benefit : fast and easy to implement**

# Graph Neural Networks



## Pytorch Geometric

PyG (PyTorch Geometric) is a library built upon PyTorch to easily write and train Graph Neural Networks (GNNs) for a wide range of applications related to structured data.



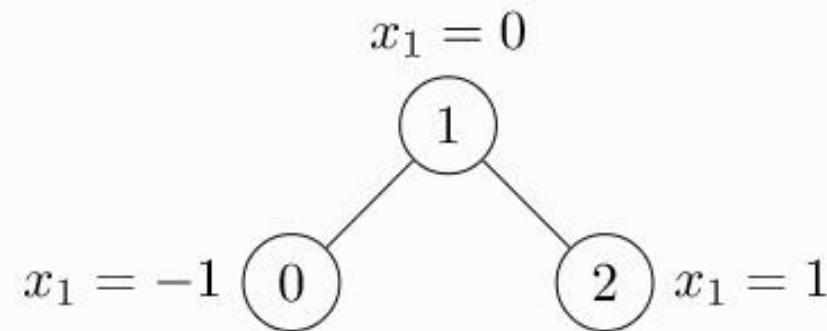
[https://pytorch-geometric.readthedocs.io/en/latest/get\\_started/colabs.html](https://pytorch-geometric.readthedocs.io/en/latest/get_started/colabs.html)

## Pytorch Geometric

```
import torch
from torch_geometric.data import Data

edge_index = torch.tensor([[0, 1, 1, 2],
                          [1, 0, 2, 1]], dtype=torch.long)
x = torch.tensor([[-1], [0], [1]], dtype=torch.float)

data = Data(x=x, edge_index=edge_index) # Graph data structure in pytorch Geometric
>>> Data(edge_index=[2, 4], x=[3, 1]) # Input data on the graph
```



Although the graph has only two edges, we need to define four index tuples to account for both directions of an edge

## Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv

conv1 = GCNConv(num_of_input_features, num_of_out_features) # init of the GNN
```

$$h_v^{k+1} = \sigma\left(W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|}\right)$$

# Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv

conv1 = GCNConv(num_of_input_features, num_of_out_features) # init of the GNN
```

$$h_v^{k+1} = \sigma(W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|})$$

```
h = conv1(input_features, graph_edges) # running the GNN on an input
```

## Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv

conv1 = GCNConv(num_of_input_features, num_of_out_features) # init of the GNN
```

$$h_v^{k+1} = \sigma(W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|})$$

```
h = conv1(input_features, graph_edges) # running the GNN on an input
```

# Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv

class GCN(torch.nn.Module):
    def __init__(self):
        super().__init__()
        torch.manual_seed(1234)
        self.conv1 = GCNConv(dataset.num_features, 4)
        self.conv2 = GCNConv(4, 4)
        self.conv3 = GCNConv(4, 2)
        self.classifier = Linear(2, dataset.num_classes)

    def forward(self, x, edge_index):
        h = self.conv1(x, edge_index)
        h = h.tanh()
        h = self.conv2(h, edge_index)
        h = h.tanh()
        h = self.conv3(h, edge_index)
        h = h.tanh() # Final GNN embedding space.

        # Apply a final (linear) classifier.
        out = self.classifier(h)

    return out, h
```

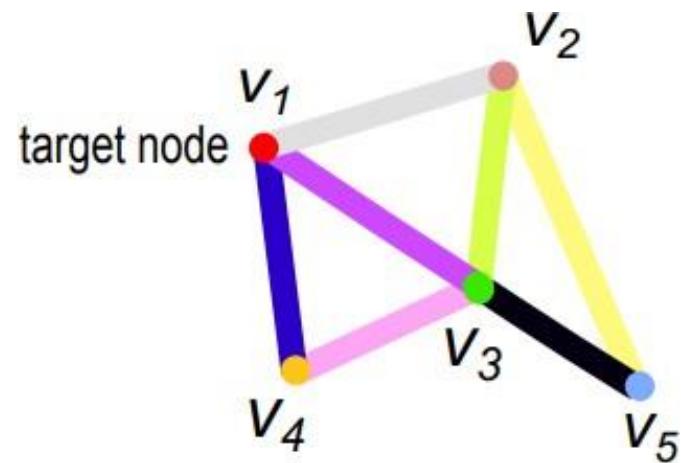
## Pytorch Geometric

```
model = GCN()
criterion = torch.nn.CrossEntropyLoss() # Define loss criterion.
optimizer = torch.optim.Adam(model.parameters(), lr=0.01) # Define optimizer.

def train(data):
    optimizer.zero_grad() # Clear gradients.
    out, h = model(data.x, data.edge_index) # Perform a single forward pass.
    loss = criterion(out[data.train_mask], data.y[data.train_mask]) # Compute the loss solely based on the training nodes.
    loss.backward() # Derive gradients.
    optimizer.step() # Update parameters based on gradients.
    return loss, h
```

# Message Passing Neural Networks

1. A graph  $G = (V, E)$ .
2. For each node  $i \in V$  we have an initial vector  $h_i^{(0)} \in \mathbb{R}^{l_0}$ .

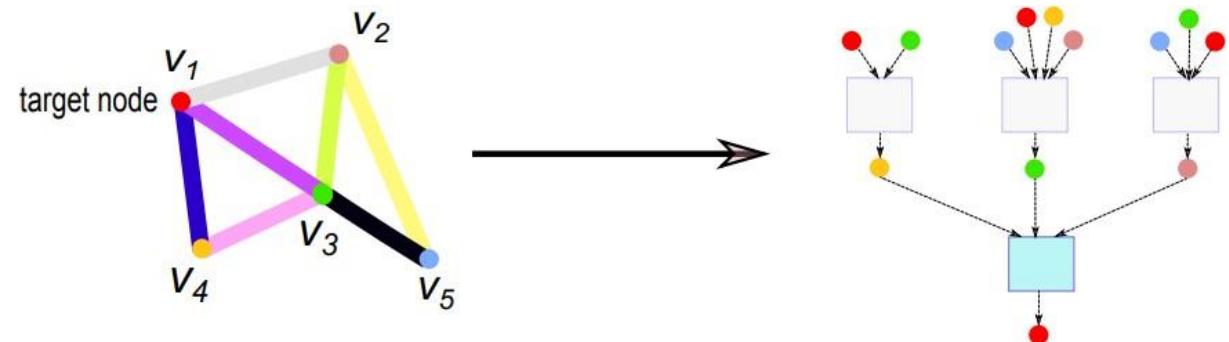


# Message Passing Neural Networks

1. A graph  $G = (V, E)$ .
2. For each node  $i \in V$  we have an initial vector  $h_i^{(0)} \in \mathbb{R}^{l_0}$ .

Given the above data, the feedforward neural algorithm on  $G$  executes  $L$  message passing schemes defined recursively for  $0 \leq k \leq L$  by:

$$h_i^{(k)} := \alpha^k \left( h_i^{(k-1)}, E_{j \in \mathcal{N}(i)} \left( \phi^k(h_i^{(k-1)}, h_j^{(k-1)}, e_{i,j}) \right) \right) \in \mathbb{R}^{l_k},$$



# Message Passing Neural Networks

1. A graph  $G = (V, E)$ .
2. For each node  $i \in V$  we have an initial vector  $h_i^{(0)} \in \mathbb{R}^{l_0}$ .

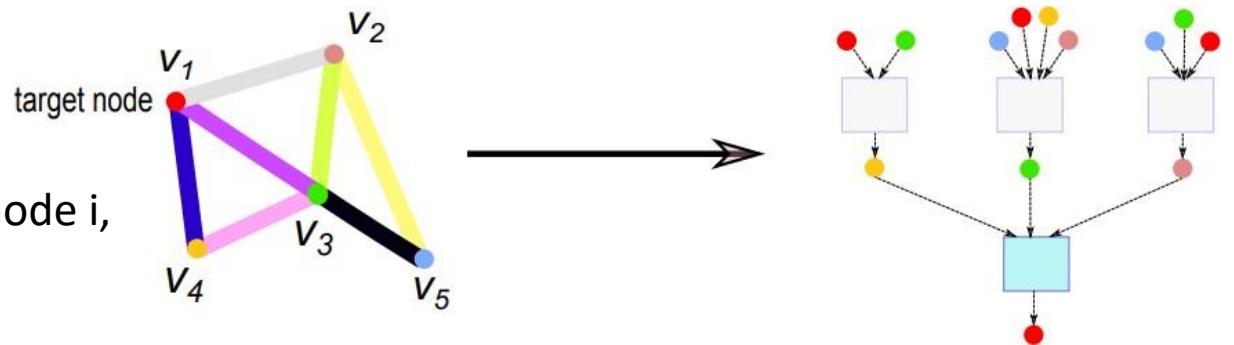
Given the above data, the feedforward neural algorithm on  $G$  executes  $L$  message passing schemes defined recursively for  $0 \leq k \leq L$  by:

$$h_i^{(k)} := \alpha^k \left( h_i^{(k-1)}, E_{j \in \mathcal{N}(i)} \left( \phi^k(h_i^{(k-1)}, h_j^{(k-1)}, e_{i,j}) \right) \right) \in \mathbb{R}^{l_k},$$

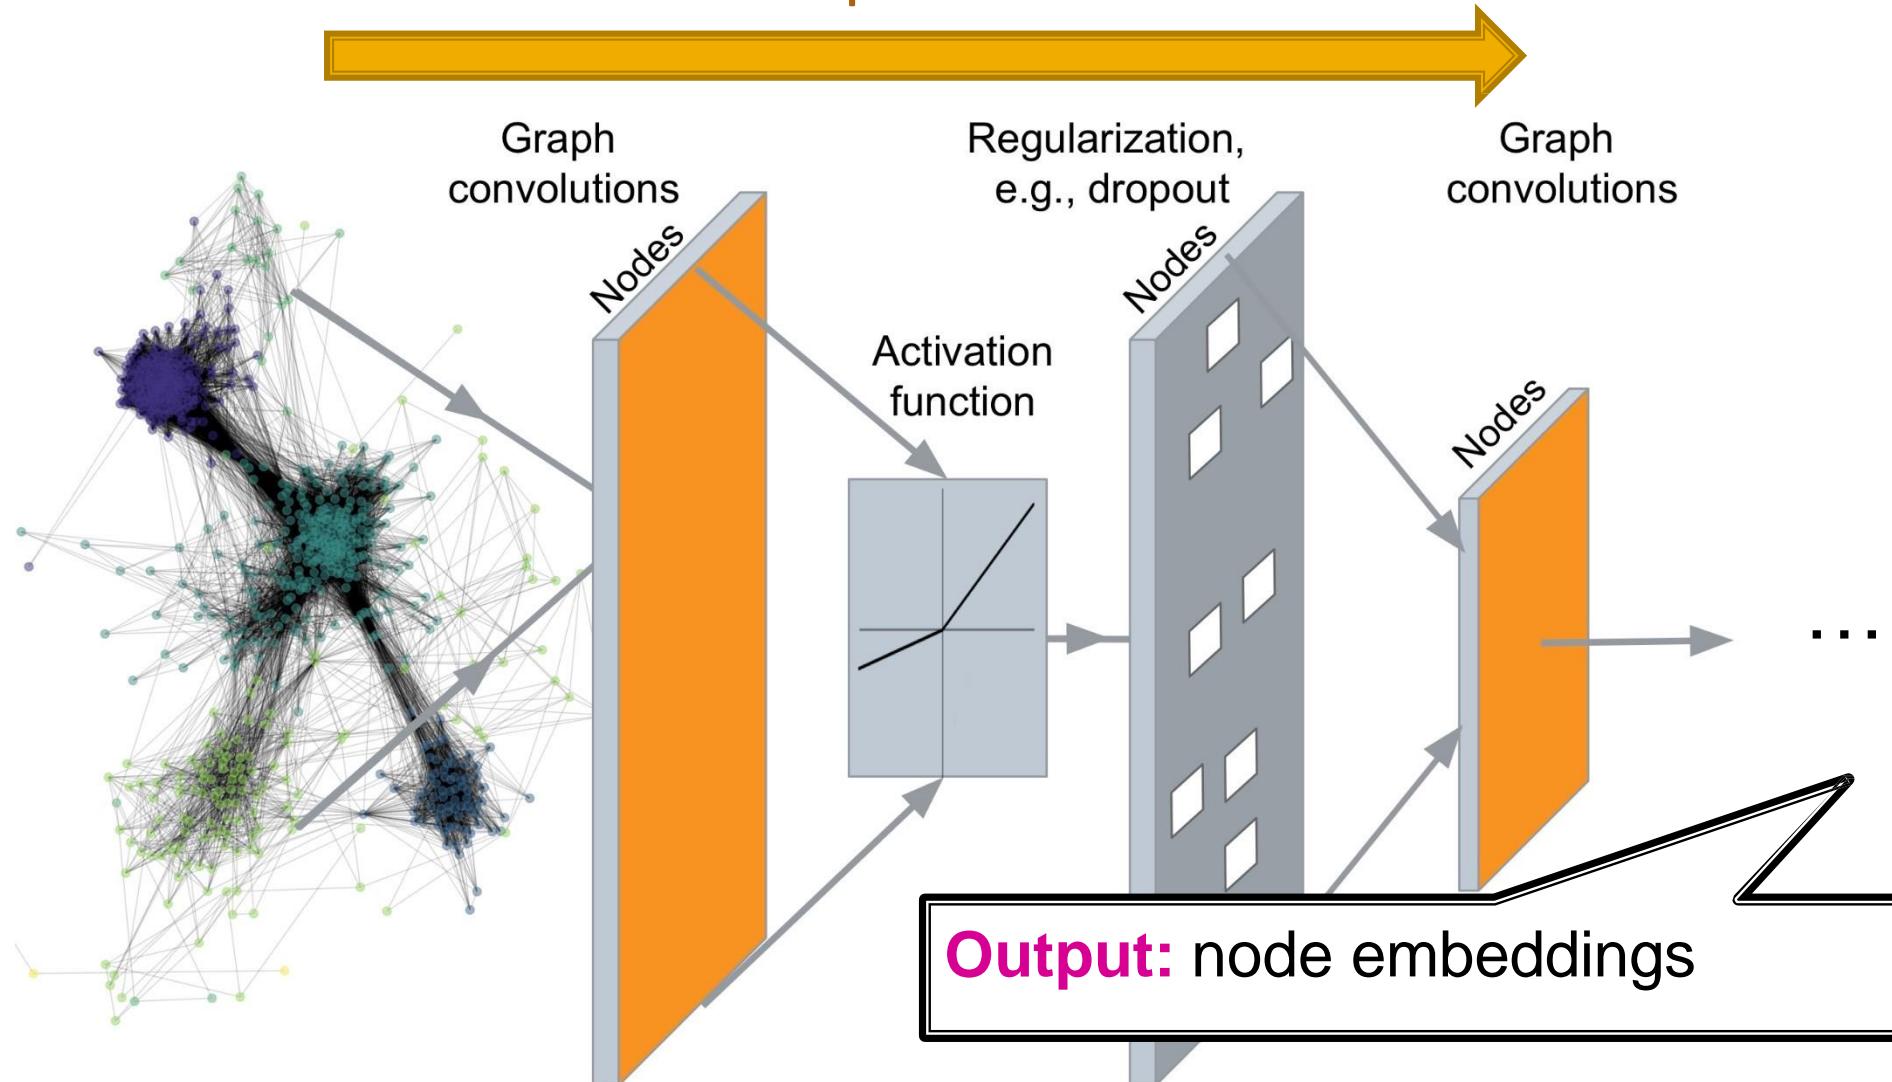
where  $e_{ij} \in R^D$  is an edge feature from the node  $j$  to the node  $i$ ,

$E$  is a permutation invariant differentiable function,

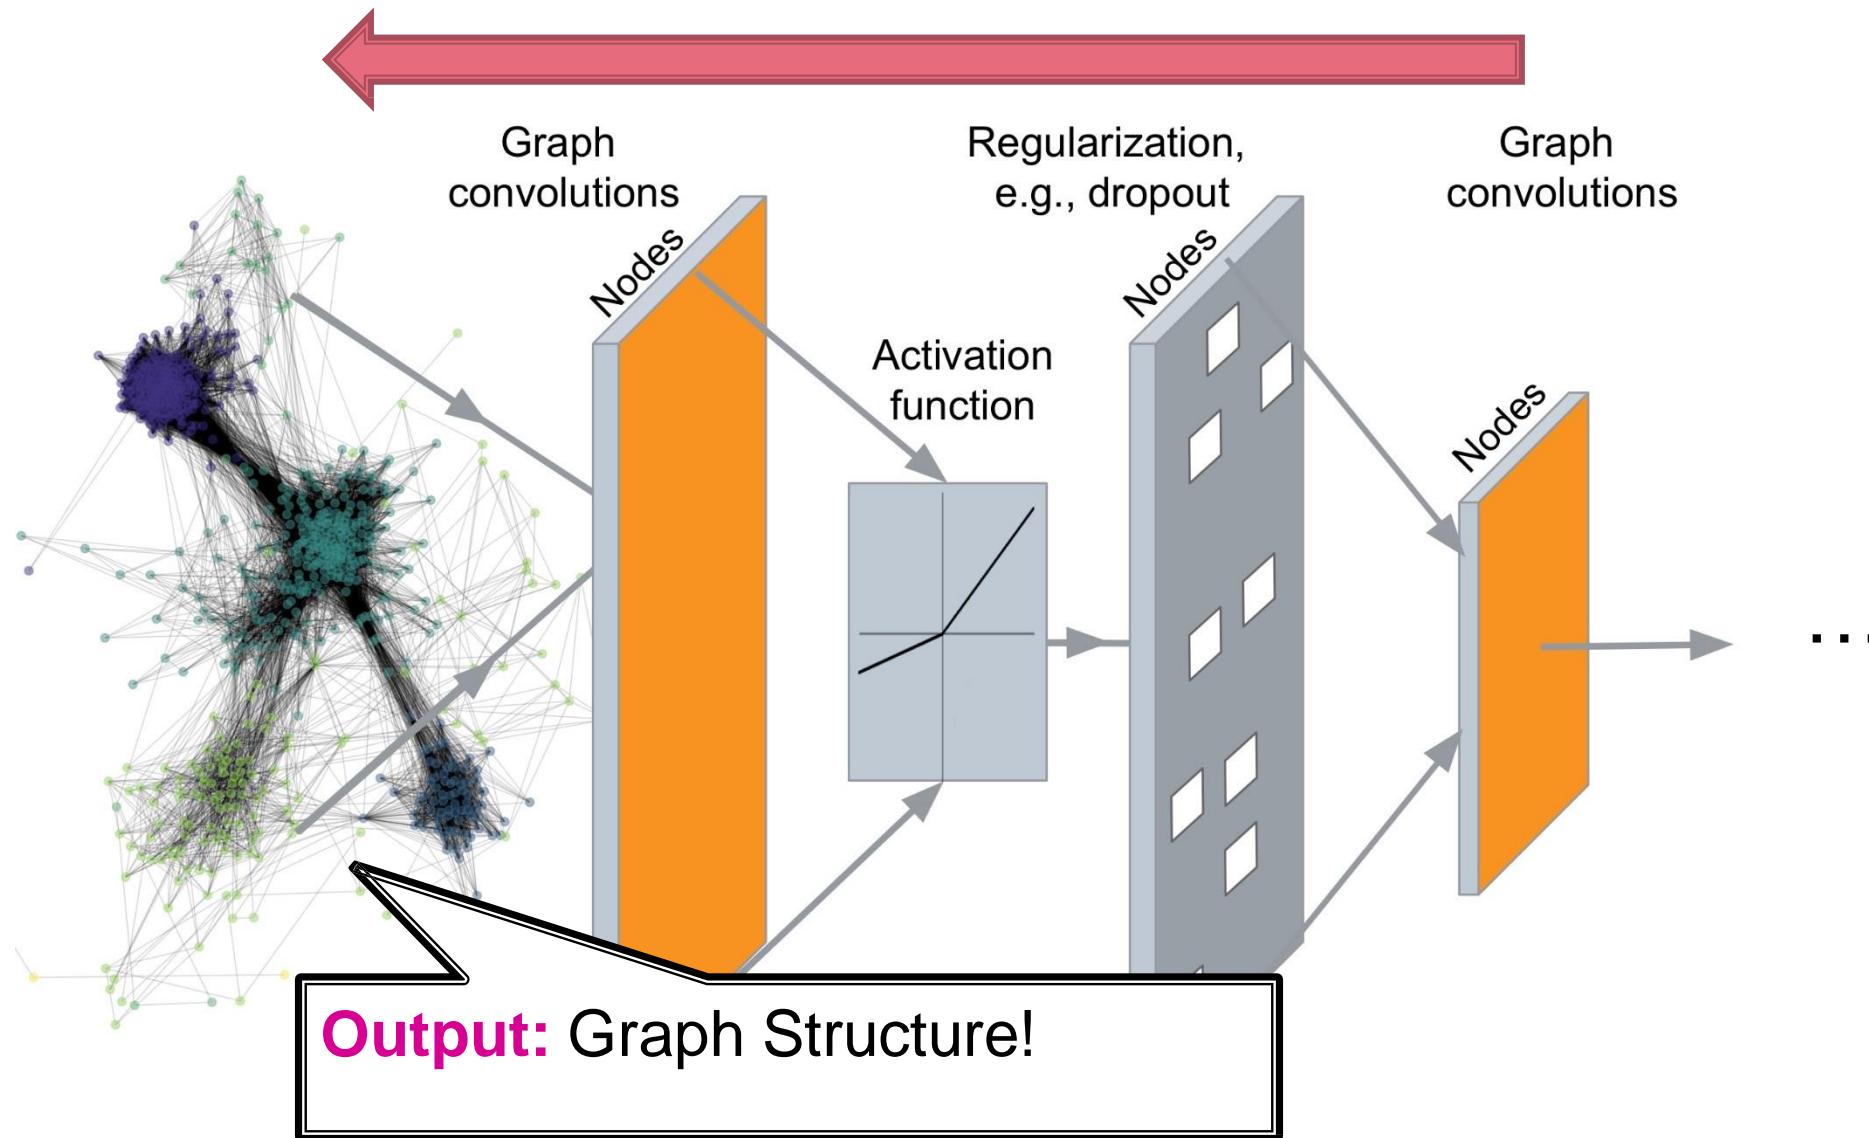
$\alpha^k, \phi^k$  are trainable differentiable functions (regular neural networks)



# Graph Encoder



# Graph Decoder



## Graph Generation task

### **Task 1: Realistic graph generation**

- Generate graphs that are **similar to a given set of graphs**
- **Task 2: Goal-directed graph generation**
- Generate graphs that **optimize given objectives/constraints**
  - E.g., Drug molecule generation/optimization

# Graph Generative Models

- **Given:** Graphs sampled from

$$p_{data}(G)$$

- **Goal:**

- Learn the distribution  $p_{model}(G)$
- Sample from  $p_{model}(G)$

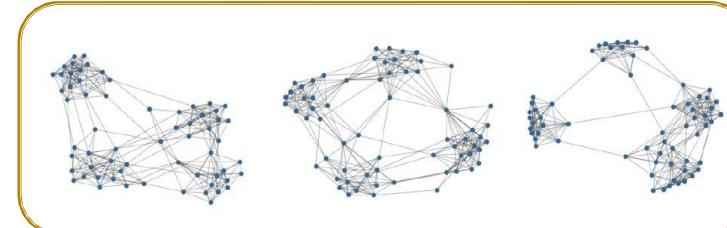
$$p_{data}(G)$$



Learn &  
Sample



$$p_{model}(G)$$



## Basics

- Assume we want to learn a generative model from a set of data points (i.e., graphs)  $\{\mathbf{x}_i\}$ 
  - $p_{data}(\mathbf{x})$  is the **data distribution**, which is never known to us, but we have sampled  $\mathbf{x}_i \sim p_{data}(\mathbf{x})$
  - $p_{model}(\mathbf{x}; \theta)$  is the **model**, parametrized by  $\theta$ , that we use to approximate  $p_{data}(\mathbf{x})$
- **Goal:**
  - (1) Make  $p_{model}(\mathbf{x}; \theta)$  close to  $p_{data}(\mathbf{x})$  (**Density estimation**)
  - (2) Make sure we can sample from  $p_{model}(\mathbf{x}; \theta)$  (**Sampling**)
    - We need to generate examples (graphs) from  $p_{model}(\mathbf{x}; \theta)$

### (1) Make $p_{model}(x; \theta)$ close to $p_{data}(x)$

- Key Principle: Maximum Likelihood
- Fundamental approach to modeling distributions

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{x \sim p_{data}} \log p_{model}(x \mid \theta)$$

- Find parameters  $\theta^*$ , such that for observed data points  $x_i \sim p_{data}$  the  $\sum_i \log p_{model}(x_i; \theta^*)$  has the highest value, among all possible choices of  $\theta$ 
  - That is, find the model that is most likely to have generated the observed data  $x$

## Non-probabilistic graph auto-encoder (GAE) model

we calculate embeddings  $Z$  and the reconstructed adjacency matrix  $\hat{\mathbf{A}}$  as follows:

$$\hat{\mathbf{A}} = \sigma(\mathbf{Z}\mathbf{Z}^\top), \text{ with } \mathbf{Z} = \text{GCN}(\mathbf{X}, \mathbf{A})$$

# Probabilistic graph auto-encoder (GAE) model

The inference model

$$q(\mathbf{Z} | \mathbf{X}, \mathbf{A}) = \prod_{i=1}^N q(\mathbf{z}_i | \mathbf{X}, \mathbf{A}), \text{ with } q(\mathbf{z}_i | \mathbf{X}, \mathbf{A}) = \mathcal{N}(\mathbf{z}_i | \boldsymbol{\mu}_i, \text{diag}(\boldsymbol{\sigma}_i^2)).$$

Here

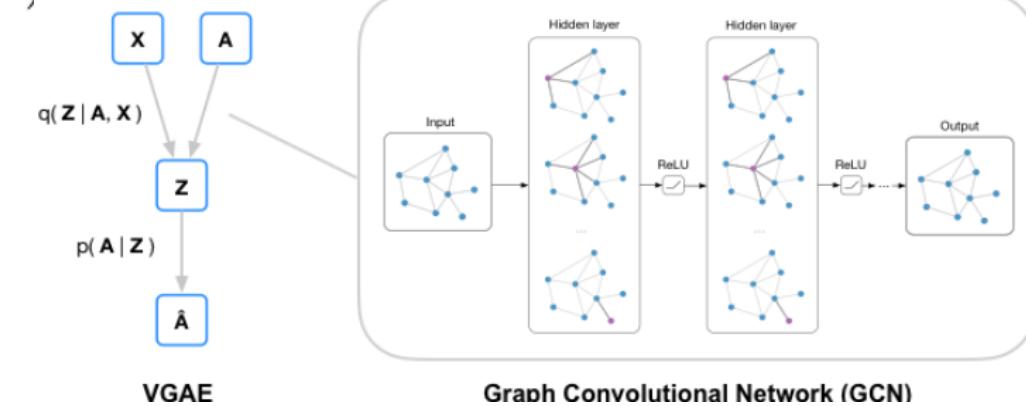
$$\boldsymbol{\mu} = \text{GCN}_{\boldsymbol{\mu}}(\mathbf{X}, \mathbf{A}) \quad \log \boldsymbol{\sigma} = \text{GCN}_{\boldsymbol{\sigma}}(\mathbf{X}, \mathbf{A})$$

Where

$$\text{GCN}(\mathbf{X}, \mathbf{A}) = \tilde{\mathbf{A}} \text{ReLU}(\tilde{\mathbf{A}} \mathbf{X} \mathbf{W}_0) \mathbf{W}_1,$$

Here also recall :

$\tilde{\mathbf{A}} = \mathbf{D}^{-\frac{1}{2}} \mathbf{A} \mathbf{D}^{-\frac{1}{2}}$  is the symmetrically normalized adjacency matrix.



## Probabilistic graph auto-encoder (GAE) model

Generative model

$$p(\mathbf{A} | \mathbf{Z}) = \prod_{i=1}^N \prod_{j=1}^N p(A_{ij} | \mathbf{z}_i, \mathbf{z}_j), \text{ with } p(A_{ij} = 1 | \mathbf{z}_i, \mathbf{z}_j) = \sigma(\mathbf{z}_i^\top \mathbf{z}_j)$$

## Probabilistic graph auto-encoder (GAE) model

Generative model

$$p(\mathbf{A} | \mathbf{Z}) = \prod_{i=1}^N \prod_{j=1}^N p(A_{ij} | \mathbf{z}_i, \mathbf{z}_j), \text{ with } p(A_{ij} = 1 | \mathbf{z}_i, \mathbf{z}_j) = \sigma(\mathbf{z}_i^\top \mathbf{z}_j)$$

## Refs

[12-deep-generation.pdf \(stanford.edu\)](#)